Word Embeddings

Michael Szczepaniak

A bit about me...

- 6 years as Sr. Data Scientist
  - operations research for defense contractor
  - top secret clearance
- 7 years as software engineer
  - couple of small companies
- Last two degrees
  - MS Data Science, RIT (2024)
    - Capstone: LLM assisted model development [1, 2]
  - BS Computer Science, CSU

RIT Rochester Institute of Technology

COLORADO STATE UNIVERSITY

What we'll cover

- A few important NLP tasks
- Getting ML models to understand text
  - text processing
  - one hot encoding vs. word embeddings
- Coding examples
  - word similarity
  - relationships between words
  - logistic regression
- Summary

A few important NLP tasks

- Text classification
- Part of speech (PoS) tagging
- Named entity recognition (NER)

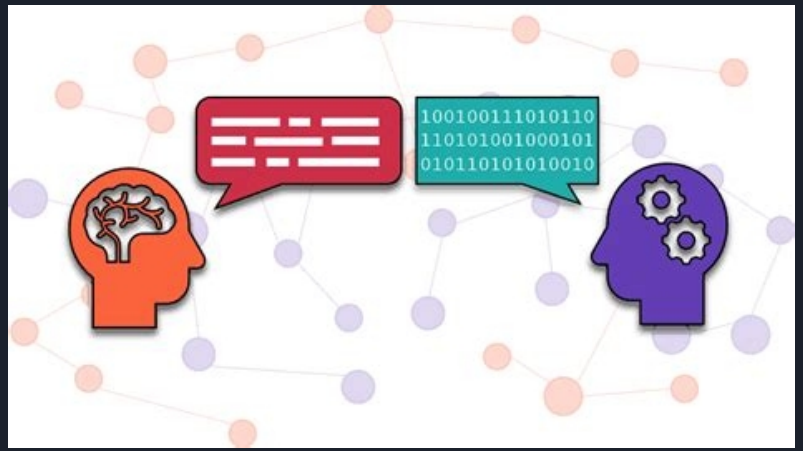

Text classification

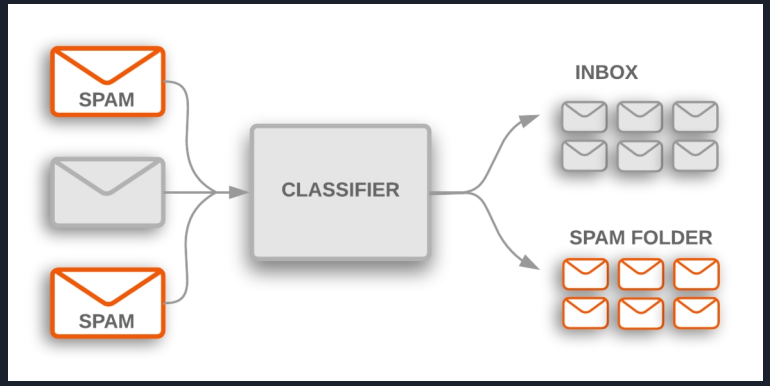

- What bucket do we put a document into?
  - binary: spam or not spam
  - multi-class: tech, sports or fashion
- Lots of models
  - Logistic regression
    - Why try this first?
  - Tree-base methods
  - Neural networks

Part of Speach (PoS) tagging

- Identify what part of speech each word is in a sentence or document.
- First step in word sense disambiguation
  - WSD arises from words having more than one meaning:

I need to get to the bank to deposit my canoe.

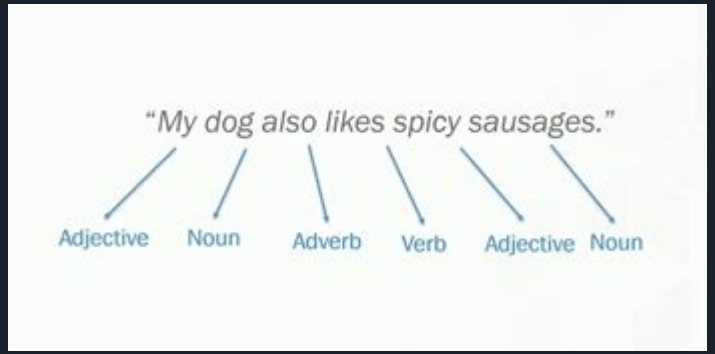

Named Entity Recognition (NER)

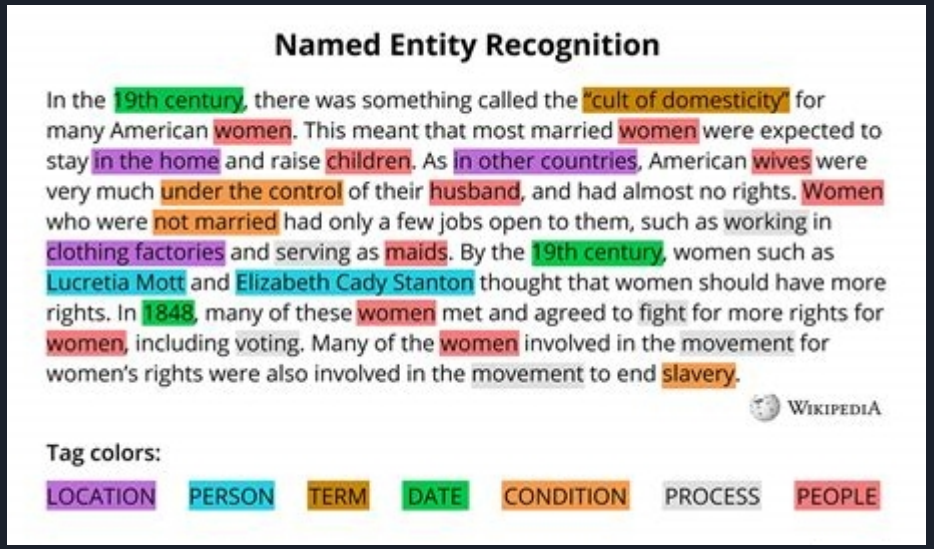

- Categorize specific types of words
  - person, organization, place/location,...
- Use cases
  - id customer names in customer service transcript
  - determine location of a user in a social media post
  - content discovery

Getting ML models to understand text

- Why we need encoding
  - models built for numeric inputs
- Simplest way to convert text to numbers
  - treat each word as a categorical variable: one hot encoding (vector of integers)
- Pros of one hot encoding
  - simple word encoding
  - simple document encoding
- Cons of one hot encoding
  - words have no relationship to each other
  - word vector size = | vocabulary | (27 - 57k words [5])

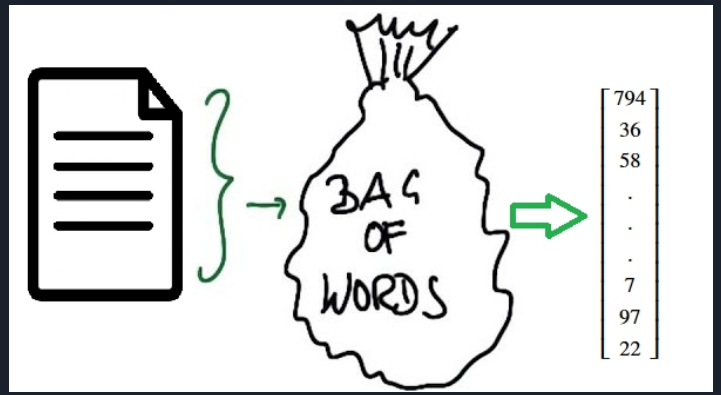

Getting ML models to understand text (cont.)

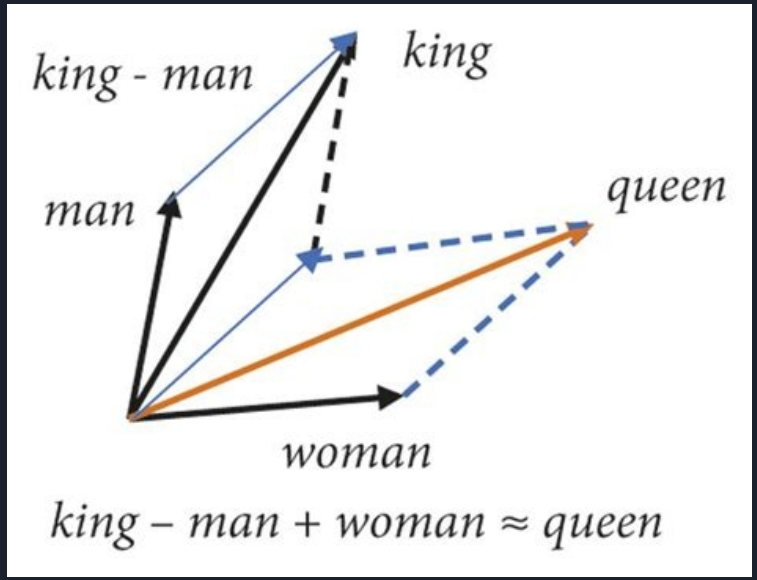

- Word embeddings capture relationships
  - vectors of real number (not integers)
- Number of dimensions: meta-parameter
  - performance vs. resolution trade-off
- Pros of word embeddings
  - captures relationship between words
  - higher information density
- Cons of word embeddings
  - must choose: canned vs. custom
  - must choose algorithm (e.g. GloVe, Word2vec, et. al.)

Coding examples

This section done in a python jupyter notebook

Summary

- A few of the important NLP tasks
  - text classification
  - PoS tagging
  - NER
  - WSD
- NLP ML model need numeric input
  - one hot encoding (simplest)
  - word embeddings (pre-trained /canned or custom)
  - pros and cons: one hot encoding vs. word embeddings

Summary (cont.)

- Coding examples: word similarity & relationships
  - words with similar meaning closest together in embedding space
  - ([king] - [man]) more similar to ([queen] - [woman]) than ([cat] - [dog])
- Coding example: logistic regression
  - Doing logistic regression as a first model is a good idea
  - text processing
  - word encoding → document encoding
  - information density: word embeddings > one hot encoding
- Custom or canned embeddings?
  - Cost versus performance trade-off

Appendices

Appendix A - What are word embeddings?

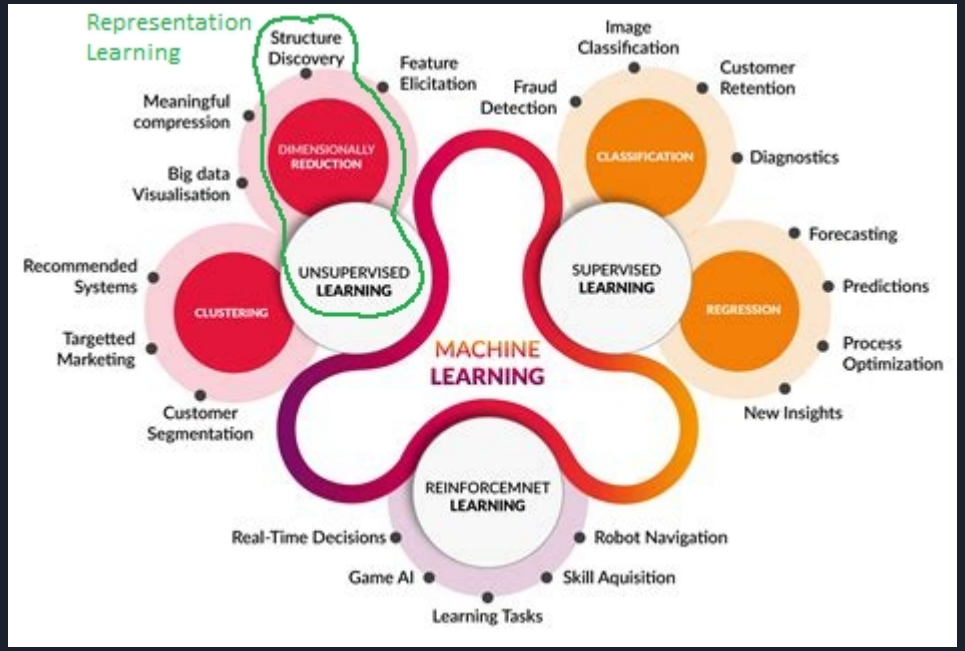

- Mapping of words to vectors
- Results from representation learning
  - Unsupervised learning
- More than one way to create
  - GloVe (Stanford)
  - Word2Vec (Google)
  - Canned vs. Custom
  - Custom: must select algorithm

Appendix B - Text processing example

I can't bloody wait!! Sony Sets a Date For Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª #stephenking #thedarktower http://t.co/J9LPdRXCDE @bdisgusting

➔ **normalize URLs**

I can't bloody wait!! Sony Sets a Date For Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª #stephenking #thedarktower <url> @bdisgusting

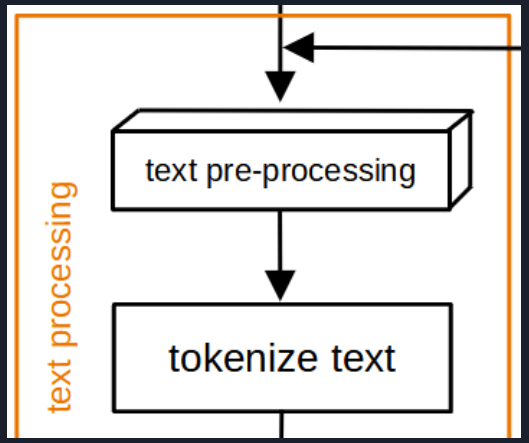

➔ **normalize twitter special chars**

I can't bloody wait!! Sony Sets a Date For Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª <hashtag>stephenking <hashtag>thedarktower <url> <user>bdisgusting

➔ **expand contractions**

I can not bloody wait!! Sony Sets a Date For Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª <hashtag>stephenking <hashtag>thedarktower <url> <user>bdisgusting

➔ **remove stop words**

~~I can~~ not bloody wait!! Sony Sets ~~a~~ Date ~~For~~ Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª <hashtag>stephenking <hashtag>thedarktower <url> <user>bdisgusting

➔ **remove punc, lemmatize, lower case, remove singletons and OOV words**

not bloody wait!! Sony Sets Date Stephen King⑨Ûªs ⑨Û÷The Dark Tower⑨Ûª <hashtag>stephenking <hashtag>thedarktower <url> <user>bdisgusting

not bloody wait sony set date stephen dark <hashtag> <hashtag> <url> <user>

➔ **tokenize**

not | bloody | wait | sony | set | date | stephen | dark | <hashtag> | <hashtag> | <url> | <user>

References

1. Large Language Model Assisted Model Development - Final paper
   https://github.com/MichaelSzczepaniak/llmamd/blob/main/docs/Final_paper.pdf

2. Large Language Model Assisted Model Development project code, data, notebooks
   https://github.com/MichaelSzczepaniak/llmamd

3. Word Embeddings notebook and supporting materials
   https://github.com/MichaelSzczepaniak/WordEmbeddings/

4. What is Named Entity Recognition (NER): Benefits, Use Cases
   https://www.expressanalytics.com/blog/what-is-named-entity-recognition-ner-benefits-use-cases-algorithms/

5. How many words do we know?
   https://doi.org/10.3389/fpsyg.2016.01116